

Design Document

Compute

Would you run it as Lambda or containers (ECS/Fargate)? Why?

Option A: AWS Lambda (Serverless)

- **Pros:**
 - No server management
 - Auto-scaling
 - Pay per request
 - Good for variable traffic
- **Cons:**
 - Cold starts
 - 15-minute timeout limit
 - Limited control over runtime
- **Use Case:** Low to moderate traffic, event-driven processing

Option B: ECS/Fargate (Containers)

- **Pros:**
 - More control
 - No cold starts
 - Long-running processes
 - Better for high throughput
- **Cons:**
 - More operational overhead
 - Cost even when idle
- **Use Case:** High traffic, predictable load, need for long-running processes

Recommendation: Use **Lambda** for feed handlers and API endpoints (stateless, request-driven). Use **ECS/Fargate** for order/error processors (long-running, continuous processing).

Streaming

What would replace `valid_orders` / `error_orders` (SQS, EventBridge, Kinesis, MSK)?

Current Implementation: In-memory queues with subscriber pattern.

Production Design Options:

Option A: Amazon SQS

- **Pros:** Simple, managed, built-in retry, dead-letter queues
- **Cons:** Not ideal for high-throughput, limited ordering guarantees
- **Use Case:** Good for moderate volume (< 1000 messages/sec per partner)

Option B: Amazon EventBridge

- **Pros:** Event-driven, rule-based routing, integration with many AWS services
- **Cons:** Less control over message ordering
- **Use Case:** Good for event-driven architecture with multiple consumers

Option C: Amazon Kinesis Data Streams

- **Pros:** High throughput, ordered processing, replay capability
- **Cons:** More complex, higher cost, requires shard management
- **Use Case:** High-volume, real-time processing needs

Option D: Amazon MSK (Managed Kafka)

- **Pros:** Industry standard, high throughput, strong ordering, multiple consumers
- **Cons:** Most complex, higher operational overhead
- **Use Case:** Enterprise-grade, high-volume, multiple consumer groups

Recommendation: Start with **SQS** for simplicity and cost-effectiveness. Migrate to **Kinesis** or **MSK** if volume exceeds SQS limits or ordering becomes critical.

Storage

What would you use for orders and errors (DynamoDB vs RDS)? Why?

Current Implementation: In-memory Map and arrays.

Production Design Options:

Option A: Amazon DynamoDB

- **Pros:**
 - Serverless, auto-scaling
 - Fast queries with GSI
 - Built-in TTL for error cleanup
 - Strong consistency for sequence numbers
- **Cons:**
 - Query patterns limited by key design
 - Cost at scale
- **Use Case:** High-throughput writes, simple queries, serverless architecture

Option B: Amazon RDS (PostgreSQL/MySQL)

- **Pros:**
 - SQL queries, complex aggregations
 - ACID transactions
 - Familiar for most developers
- **Cons:**
 - Requires capacity planning
 - Scaling is more complex
- **Use Case:** Complex queries, relational data, existing SQL expertise

Recommendation: Use **DynamoDB** for orders (high write volume, simple queries) and **RDS** if complex analytics are needed. For errors, use **DynamoDB** with TTL.

Table Design (DynamoDB):

- **Orders Table:**
 - Partition Key: partnerId
 - Sort Key: sequenceNumber
 - GSI: eventTime-index (for time-range queries)
- **Errors Table:**
 - Partition Key: partnerId
 - Sort Key: receivedTime
 - TTL: expirationTime (30 days)

Frontend hosting

How to host the React app (S3 + CloudFront recommended)

Current Implementation: React app served by Vite dev server.

Production Design:

- **Hosting:** S3 + CloudFront
 - Build React app: `npm run build`
 - Upload `dist/` to S3 bucket

- Configure CloudFront distribution
- Custom domain with Route 53
- **API Integration:**
 - API Gateway endpoint
 - CORS configured
 - Environment variables for API URL

Configuration & secrets

env vars, parameter store, secrets manager

Environment Variables

- Store in **AWS Systems Manager Parameter Store** (for non-sensitive config)
- Store in **AWS Secrets Manager** (for API keys, database credentials)
- Use IAM roles for service-to-service authentication

Configuration Examples

- Partner configurations (API keys, endpoints)
- Validation rules
- Stream names
- Database connection strings

Observability

logs/metrics/alarms you'd set up

Logging

- **CloudWatch Logs:** Centralized logging for all services
- Structured logging (JSON format)
- Log retention: 30 days (configurable)

Metrics

- **CloudWatch Metrics:**
 - Request count, latency (API Gateway)
 - Processing time (Lambda/ECS)

- Queue depth (SQS)
- Error rate
- Order volume per partner
- **Custom Metrics:**
 - Orders processed per minute
 - Validation failure rate
 - Average processing time

Alarms

- High error rate (> 5% of requests)
- Queue depth threshold (SQS > 1000 messages)
- API latency > 1 second (p95)
- Processing failures
- Storage capacity warnings

Tracing

- **AWS X-Ray:** Distributed tracing for request flows
- Track requests from API Gateway through Lambda/ECS to DynamoDB

Dashboards

- **CloudWatch Dashboards:**
 - Real-time order processing rate
 - Error rate by partner
 - API performance
 - Queue metrics

CI/CD

pipeline stages (lint/test/build/package/deploy), and how you'd roll out safely

Pipeline Stages

1. **Source:** GitHub/GitLab/Bitbucket
2. **Build:**
 - Install dependencies
 - Run linters (ESLint, Prettier)
 - Run unit tests

- Build TypeScript
- 3. **Test:**
 - Integration tests
 - End-to-end tests (optional)
- 4. **Package:**
 - Create Docker images (if using ECS)
 - Package Lambda functions
- 5. **Deploy:**
 - Deploy to staging environment
 - Run smoke tests
 - Deploy to production (with approval)

Deployment Strategy

Blue/Green Deployment (ECS):

- Deploy new version alongside old
- Route traffic gradually (10% → 50% → 100%)
- Rollback by switching traffic back

Canary Deployment (Lambda):

- Deploy new version to 10% of traffic
- Monitor metrics
- Gradually increase to 100%

Infrastructure as Code:

- **AWS CDK or Terraform** for infrastructure
- Version controlled
- Repeatable deployments

Pipeline Tools

- **AWS CodePipeline + CodeBuild + CodeDeploy**
- Or **GitHub Actions / GitLab CI** with AWS deployment